

Executive classifier decision guide

State length × question count × quality × cost

Fixed labels

ModernBERT trained

lowest modeled cost

Runtime labels

GLiClass Modern

self-hosted flexibility

Managed API

JEV Noul


quality / cost balance

Mark Austin

Independent synthetic evaluation · September 20, 2026

Executive decision map: economics versus runtime flexibility


Reference workload: 6,000 state tokens × 25 questions · accuracy from the locked synthetic test



The \$1/1,000 boundary is an executive guide, not a universal purchasing threshold. Flexible systems accept new questions without retraining; trained heads do not.

Executive cost snapshots across generic state/question workloads

No transcript-duration assumption · each panel specifies state length and question count directly · lower is better



Costs are modeled per 1,000 states. Fixed encoder heads require training; shared-state zero-shot and hosted approaches permit runtime labels/questions. Hidden LLM reasoning tokens are excluded.

What should an executive choose?

Select the operating model before selecting the model

Known, stable taxonomy

ModernBERT trained heads

Best modeled economics with 99.16% test accuracy. New questions require labeled data and retraining.

Runtime labels, self-hosted

GLiClass Modern

Best open zero-shot trade-off tested: 97.90% calibrated accuracy and near-trained-encoder modeled cost.

Runtime questions, managed service

JEV Noul

99.90% calibrated accuracy with shared-state-like billing and substantially lower modeled cost than the general LLMs.

Broad reasoning or maximum test score

Luna or Sol

Near-ceiling quality, but use when semantic breadth justifies the premium. Sol's advantage over calibrated JEV was one decision in 4,050.

Executive guardrails

- Quality results come from a synthetic, explicit-evidence test and are not production guarantees.
- Cost curves are modeled from token shapes, published/API rates, and H100 throughput proxies.
- Validate with real states, real questions, target hardware, and operational error costs before procurement.